

Deep Learning for Computer Vision

Image Sampling and Interpolation

Vineeth N Balasubramanian

Department of Computer Science and Engineering
Indian Institute of Technology, Hyderabad



NPTEL
National Institute of Technology Hyderabad

NPTEL

Cost Improvement using Convolution Theorem?

Convolution Theorem

- Fourier transform of convolution of two functions is product of their Fourier transforms:

$$F[g * h] = F[g]F[h]$$

- Convolution** in spatial domain can be obtained through **multiplication** in frequency domain!

$$g * h = F^{-1}[F[g]F[h]]$$

NPTTEL

Cost Improvement using Convolution Theorem?

Convolution Theorem

- Fourier transform of convolution of two functions is product of their Fourier transforms:



$$F[g * h] = F[g]F[h]$$

- **Convolution** in spatial domain can be obtained through **multiplication** in frequency domain!


$$g * h = F^{-1}[F[g]F[h]]$$

- Image convolution needs $O(N^2 \cdot k^2)$ time, where $N \times N$ is image size, and $k \times k$ is kernel size
- By performing convolution in Fourier domain, cost is: $O(N^2)$ for a single pass over the image + cost of FFT: $O(N^2 \log N^2)$ for the image and $O(k^2 \log k^2)$ for the kernel $\approx O(N^2 \log N^2 + k^2 \log k^2)$, in total (other terms additive)



Exercise: Match spatial domain image to Fourier magnitude image





B



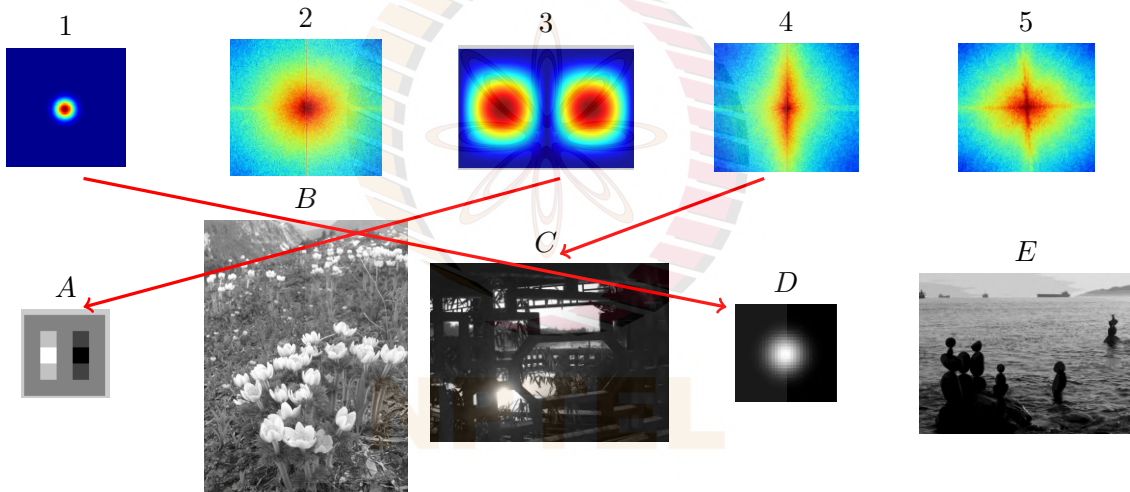
C



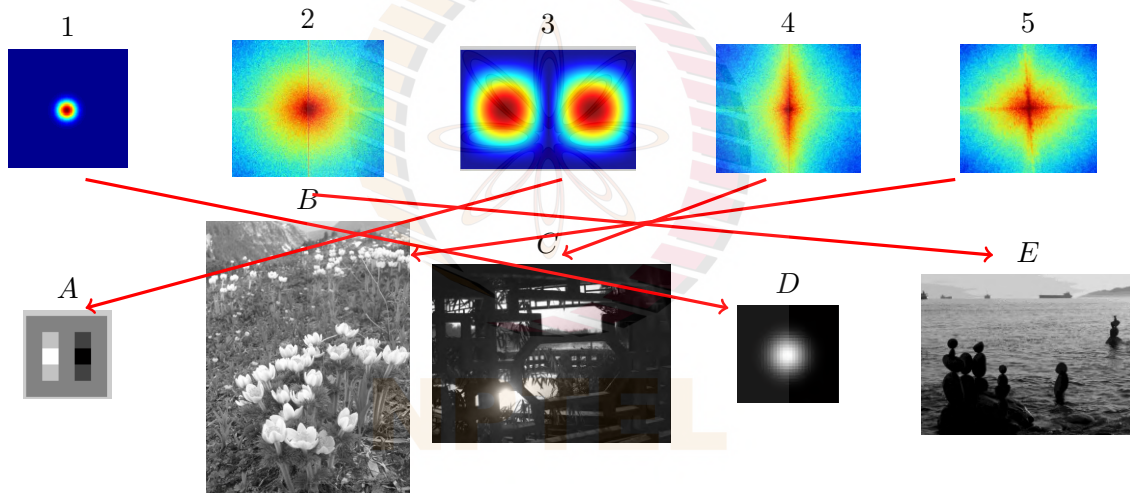
E



Exercise: Match spatial domain image to Fourier magnitude image



Exercise: Match spatial domain image to Fourier magnitude image



What sense does a low-resolution image make to us?



Original



Subsampled & zoomed

Clues from human perception:

- Early processing in human's filters for various orientations and scales of frequency.
- Perceptual cues in mid-high frequencies dominate perception.
- When we see an image from far away, we are effectively **sub-sampling** it.

Credit: Ron Hansen (Unsplash)

What sense does a low-resolution image make to us?



Original



Subsampled & zoomed


Clues from human perception:

- Early processing in human's filters for various orientations and scales of frequency.
- Perceptual cues in mid-high frequencies dominate perception.
- When we see an image from far away, we are effectively **sub-sampling** it.

Credit: Ron Hansen (Unsplash)

Sub-sampling

Throw away every other row and column to create a $1/2$ size image.



$1/4$




$1/8$

Credit: S Seitz, R Urtasun

Sub-sampling


Why does this look so cruffy?



Credit: S Seitz, R Urtasun


Sub-sampling

What's happening?



Credit: S Seitz, R Urtasun

Aliasing




- Occurs when your sampling rate is not high enough to capture the amount of detail in your image.
- To do sampling right, need to understand the structure of your signal/image.
- The minimum sampling rate is called the **Nyquist rate**.

Aliased



Aliasing




- Occurs when your sampling rate is not high enough to capture the amount of detail in your image.
- To do sampling right, need to understand the structure of your signal/image.
- The minimum sampling rate is called the **Nyquist rate**.

Aliased



Aliasing



- Occurs when your sampling rate is not high enough to capture the amount of detail in your image.
- To do sampling right, need to understand the structure of your signal/image.
- The minimum sampling rate is called the **Nyquist rate**.


Aliased



Aliasing: Problems

Shannon's Sampling Theorem shows that the minimum sampling is:

$$f_s \geq 2f_{max}$$



Example of a 1D signal


Examples

- Image
 - Striped shirt's pattern look weird on screen.
- Video
 - **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.
- Graphics
 - Checkerboards disintegrate in ray tracing.

Aliasing: Problems

Shannon's Sampling Theorem shows that the minimum sampling is:

$$f_s \geq 2f_{max}$$



Example of a 1D signal


Examples

- **Image**
 - Striped shirt's pattern look weird on screen.
- **Video**
 - **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.
- **Graphics**
 - Checkerboards disintegrate in ray tracing.

Aliasing: Problems

Shannon's Sampling Theorem shows that the minimum sampling is:

$$f_s \geq 2f_{max}$$




Examples

- **Image**
 - Striped shirt's pattern look weird on screen.
- **Video**
 - **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.
- **Graphics**
 - Checkerboards disintegrate in ray tracing.

Aliasing: Problems

Shannon's Sampling Theorem shows that the minimum sampling is:

$$f_s \geq 2f_{max}$$



Examples


- **Image**
 - Striped shirt's pattern look weird on screen.
- **Video**
 - **Wagon Wheel effect:** Wheels spins in the opposite direction at high speed.
- **Graphics**
 - Checkerboards disintegrate in ray tracing.

Aliasing: Image



Striped shirt's pattern look weird on screen.


Aliasing : Video



Wagon Wheel effect: Wheels spins in the opposite direction at high speed.

NPTTEL

Aliasing: Graphics



Checkerboards disintegrate in ray tracing.

Aliasing: Nyquist Limit 2D example




Credit: S Seitz, R Urtasun

Anti-aliasing




Anti-aliasing

Example: Gaussian Pre-filtering



Credit: N Snavely, R Urtasun



Subsampling with Gaussian Pre-filtering




Credit: N Snavely, R Urtasun

Upsampling





x



$10x$

How to go from left to right? **Interpolation**. Simple method: Repeat each row and column 10 times (Nearest Neighbour Interpolation).

Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.



What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

Interpolation



Recall how a digital image is formed,

$$F[x, y] = \text{quantize}\{f(xd, yd)\}$$

- It is a discrete point-sampling of a continuous function.
- If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale.

What if we don't know f ?

- Guess an approximation: Can be done in a principled way via filtering.
- Convert F to a continuous function:

$$f_F(x) = \begin{cases} F(\frac{x}{d}) & \text{if } \frac{x}{d} \text{ is an integer} \\ 0 & \text{otherwise} \end{cases}$$

- Reconstruct: $\hat{f} = h * f_F$

Interpolation as Convolution

- To **interpolate** (or upsample) an image to a higher resolution, we need an **interpolation kernel** with which to **convolve** the image:

$$g(i, j) = \sum_{k, l} f(k, l) h(i - rk, j - rl)$$

Above formula similar to discrete convolution^a, except that we replace k and l in $h(\cdot)$ with rk and rl , where r is the upsampling rate.

- Linear interpolator (corresponding to *tent kernel*) produces interpolating piecewise linear curves.
- More complex kernels e.g., B-splines.

^a $g = f * h \implies g(i, j) = \sum_{k, l} f(k, l) h(i - k, j - l)$

Interpolation as Convolution

- To **interpolate** (or upsample) an image to a higher resolution, we need an **interpolation kernel** with which to **convolve** the image:



$$g(i, j) = \sum_{k, l} f(k, l) h(i - rk, j - rl)$$

Above formula similar to discrete convolution^a, except that we replace k and l in $h(\cdot)$ with rk and rl , where r is the upsampling rate.


- Linear interpolator (corresponding to *tent kernel*) produces interpolating piecewise linear curves.
- More complex kernels e.g., B-splines.

^a $g = f * h \implies g(i, j) = \sum_{k, l} f(k, l) h(i - k, j - l)$


Types of Interpolation




"Ideal" reconstruction



Nearest-neighbor interpolation



Linear interpolation



Gaussian reconstruction


Credit: B Curless

Examples

Original Image:



Upsampled Images:



Left to right: Nearest Neighbour Interpolation, Bilinear Interpolation, Bicubic Interpolation.

Interpolation and Decimation

Interpolation

To **interpolate** (or upsample) an image to a higher resolution, we need an **interpolation kernel** with which to convolve the image (r is upsampling rate):

$$g(i, j) = \sum_{k, l} f(k, l) h(i - rk, j - rl)$$

Decimation (Sub-sampling)

To **decimate** (or sub-sample) an image to a lower resolution, we need an **decimation kernel** with which to convolve the image (r is downsampling rate):

$$g(i, j) = \sum_{k, l} f(k, l) h(i - \frac{k}{r}, j - \frac{l}{r})$$

Homework

Readings

- Chapter 3 (§3.5.1-3.5.2), Szeliski, *Computer Vision: Algorithms and Applications*
- Chapter 7 (§7.4), Forsyth and Ponce, *Computer Vision: A Modern Approach*

NPTEL