Deep Learning for Computer Vision

Image Representation

Vineeth N Balasubramanian

Department of Computer Science and Engineering
Indian Institute of Technology, Hyderabad

Question

On Colour

If visible light spectrum is VIBGYOR, why RGB colour representation?

Question

On Colour

If visible light spectrum is VIBGYOR, why RGB colour representation?

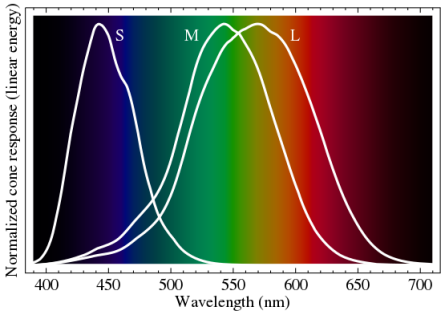

Credit: Derek Hoiem, UIUC

https://en.wikipedia.org/wiki/Color_vision

- Long (red), Medium (green), and Short (blue) cones, plus intensity rods
- Fun facts
  - "M" and "L" on the X-chromosome $\implies$ men are more likely to be colour blind!
  - Some animals have 1 (night animals), 2 (e.g., dogs), 4 (fish, birds), 5 (pigeons, some reptiles/amphibians), or even 12 (mantis shrimp) types of cones

Image as a Matrix

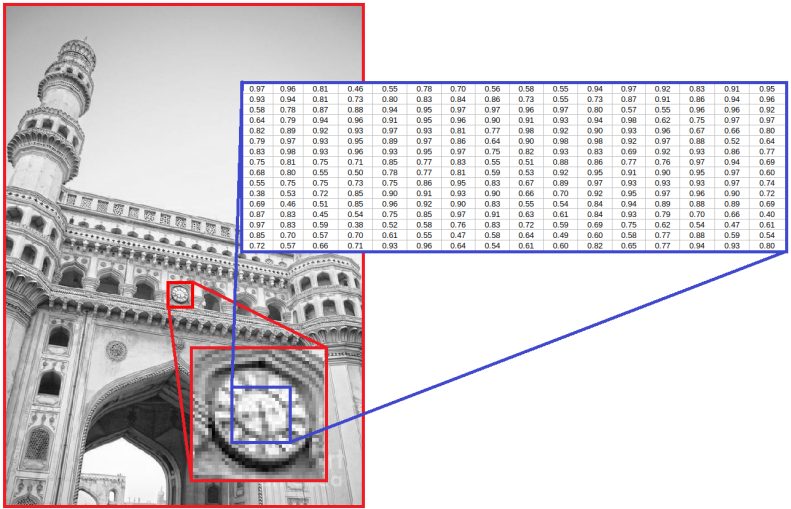

- Common to use one byte per value: 0 = black, 255 = white
- One such matrix for every channel in colour images

Image as a Function

- We can think of a (grayscale) image as a function $f : \mathbb{R}^2 \to \mathbb{R}$, giving the intensity at position (x, y)
- A digital image is a discrete (sampled, quantized) version of this function

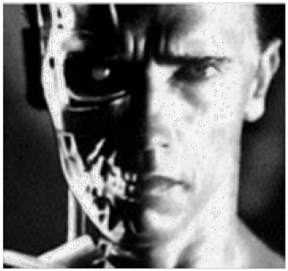

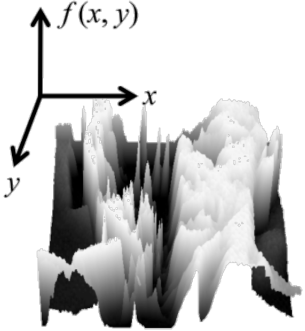

Credit: Noah Snavely, Cornell Univ

Image Transformations

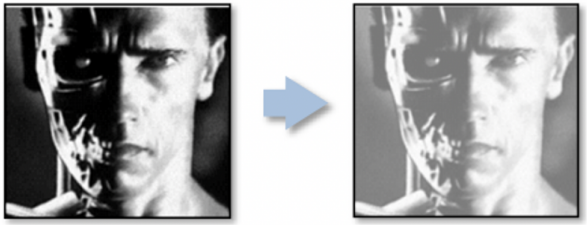

Image Transformations

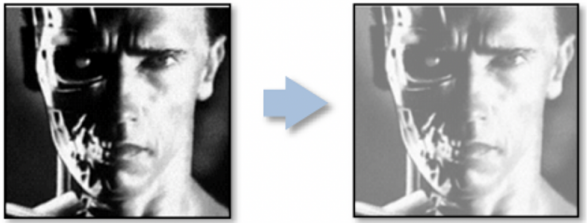

$$\hat{I}(x, y) = I(x, y) + 20$$

Image Transformations

$$\hat{I}(x, y) = I(x, y) + 20$$

Image Transformations

$$\hat{I}(x, y) = I(x, y) + 20$$

$$\hat{I}(x, y) = I(-x, y)$$

Image Processing Operations

- Point Operations
  - Output value at (m_0, n_0) is dependent only on the input value at the same coordinate
  - *Complexity/pixel:* Constant
- Local Operations
  - Output value at (m_0, n_0) is dependent on input values in a $p \times p$ neighborhood of that same coordinate
  - *Complexity/pixel:* p^2
- Global Operations
  - Output value at (m_0, n_0) is dependent on on all the values in the input $N \times N$ image
  - *Complexity/pixel:* N^2

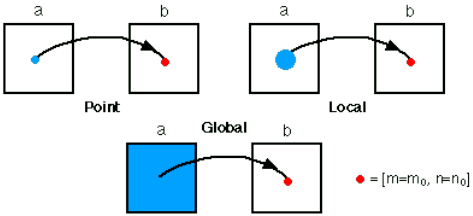

Point Operations: Example

- Image Enhancement:
  Reversing the contrast
- How?

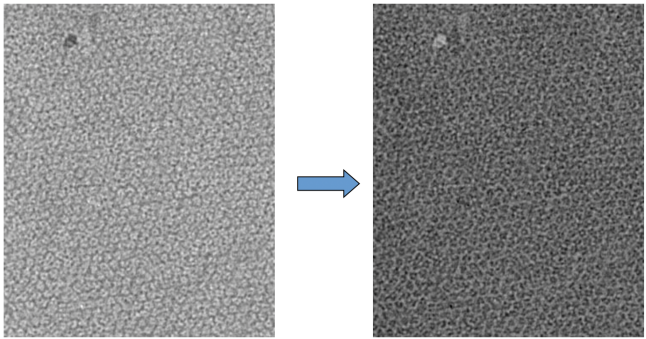

Point Operations: Example

- Image Enhancement:
  Reversing the contrast
- How?

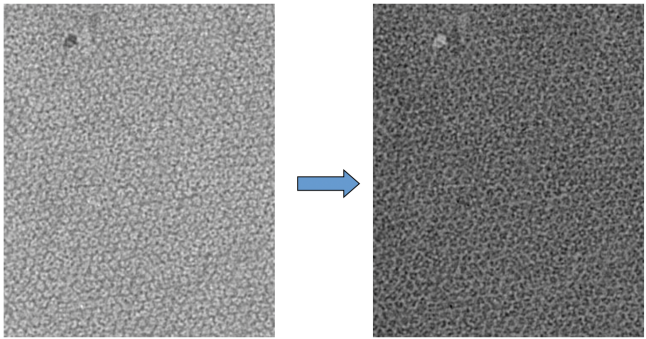

$$\hat{I}(m_0, n_0) = I_{MAX} - I(m_0, n_0) + I_{MIN}$$

Point Operations: Another Example

- Image Enhancement:
  Stretching the contrast
- How?

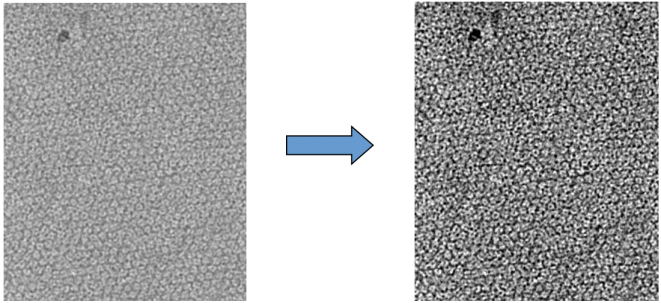

Point Operations: Another Example

- Image Enhancement:
  Stretching the contrast
- How?

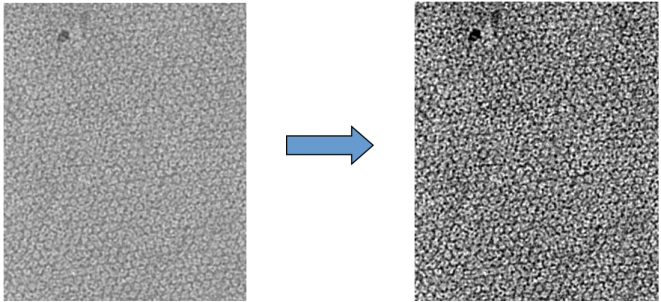

Linear Contrast Stretching

$$\hat{I}(m_0, n_0) = \left(I(m_0, n_0) - \min_{x,y} I(x, y)\right) * \left((I_{MAX} - I_{MIN})/(\max_{x,y} I(x, y) - \min_{x,y} I(x, y))\right) + I_{MIN}$$

Going Beyond Linear Stretching

Question

Heard about **Histogram Equalization**? Read about it, homework!

How Useful are Point Operations?

- A single point (or pixel)'s intensity is influenced by multiple factors, and may not tell us everything
  - Light source strength and direction
  - Surface geometry, material and nearby surfaces
  - Sensor capture properties
  - Image representation and colour
- Given a camera and a still scene, how do you reduce noise using point operations?

How Useful are Point Operations?

- A single point (or pixel)'s intensity is influenced by multiple factors, and may not tell us everything
  - Light source strength and direction
  - Surface geometry, material and nearby surfaces
  - Sensor capture properties
  - Image representation and colour
- Given a camera and a still scene, how do you reduce noise using point operations?
- Take many images, and average them!
- You need local operations otherwise. What is the local operation?

Image Processing Operations

- Point Operations
  - Output value at (m_0, n_0) is dependent only on the input value at the same coordinate
  - *Complexity/pixel:* Constant
- Local Operations
  - Output value at (m_0, n_0) is dependent on input values in a $p \times p$ neighborhood of that same coordinate
  - *Complexity/pixel:* p^2
- Global Operations
  - Output value at (m_0, n_0) is dependent on on all the values in the input $N \times N$ image
  - *Complexity/pixel:* N^2

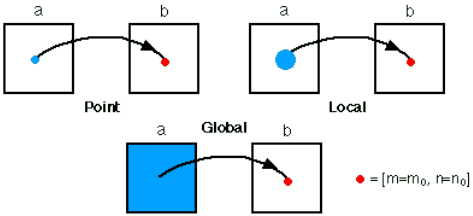

Local Operation Examples: Moving Average

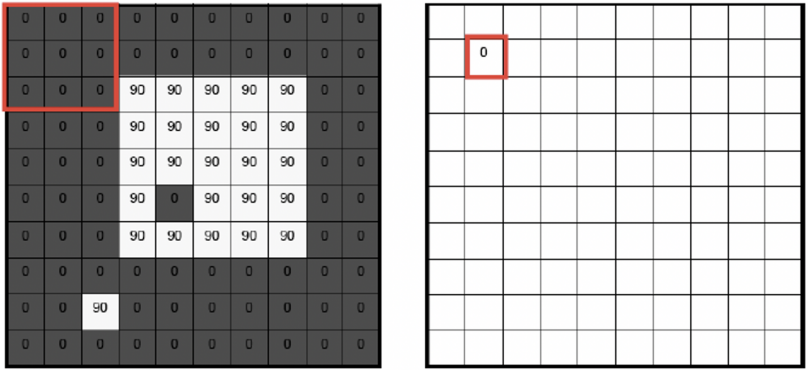

Credit: Steve Seitz, Univ of Washington

Local Operation Examples: Moving Average

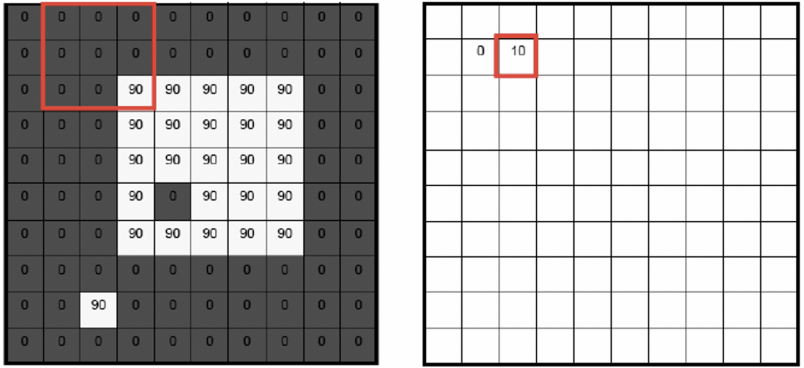

Credit: Steve Seitz, Univ of Washington

Local Operation Examples: Moving Average

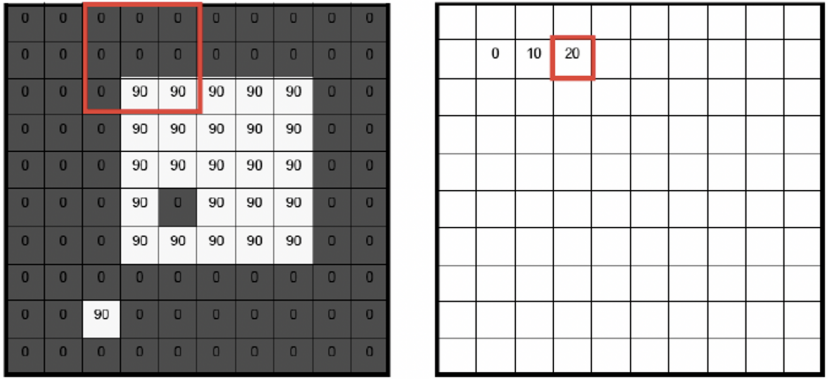

Credit: Steve Seitz, Univ of Washington

Local Operation Examples: Moving Average

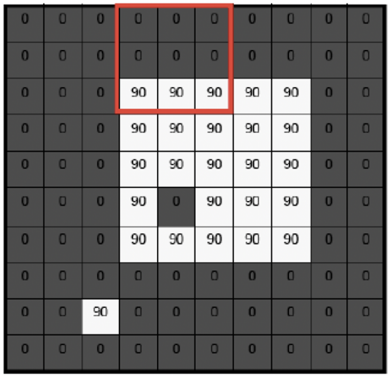

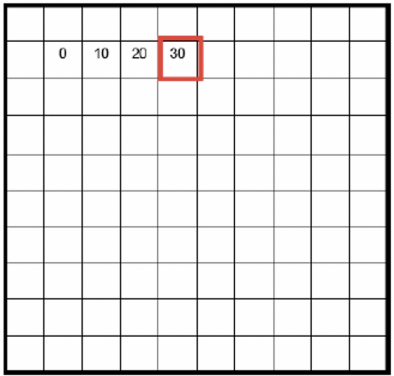

Credit: Steve Seitz, Univ of Washington

Local Operation Examples: Moving Average

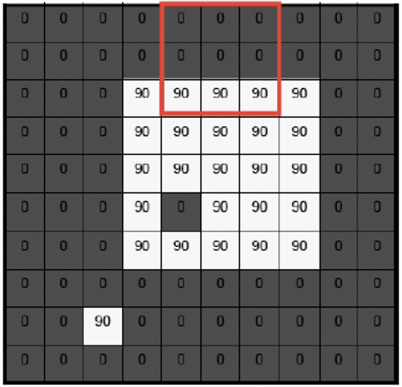

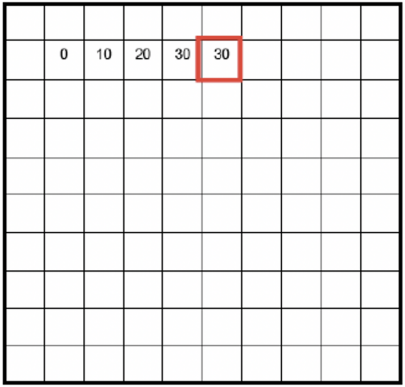

Credit: Steve Seitz, Univ of Washington

Local Operation Examples: Moving Average

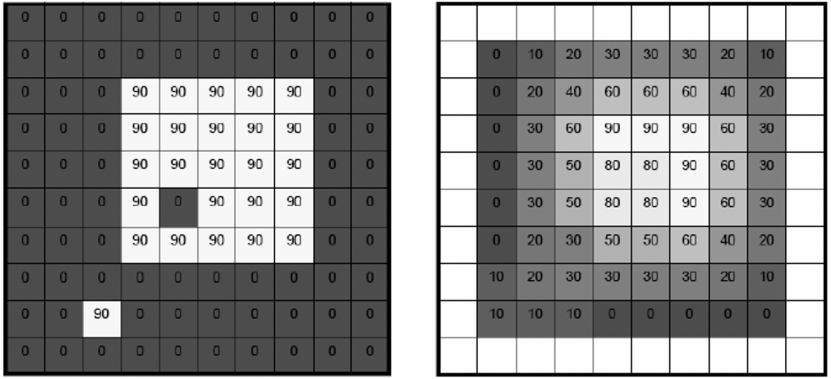

Credit: Steve Seitz, Univ of Washington

Image Processing Operations

- Point Operations
  - Output value at (m_0, n_0) is dependent only on the input value at the same coordinate
  - *Complexity/pixel:* Constant
- Local Operations
  - Output value at (m_0, n_0) is dependent on input values in a $p \times p$ neighborhood of that same coordinate
  - *Complexity/pixel:* p^2
- Global Operations
  - Output value at (m_0, n_0) is dependent on on all the values in the input $N \times N$ image
  - *Complexity/pixel:* N^2

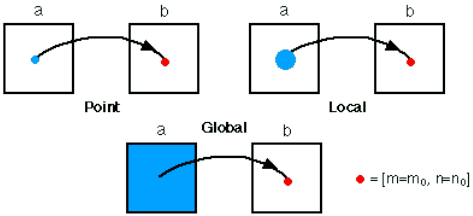

Global Operations: Examples

- Image coordinate transformations, e.g. Fourier transform
- We will see more of this later

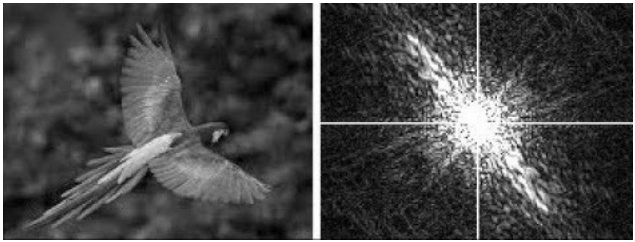

Credit: Mathworks MATLAB Toolbox

Homework

Readings

- Chapter 3.1, Szeliski, *Computer Vision: Algorithms and Applications*

Questions to Answer

- What is histogram equalization, and how do you derive its formula?

References


Richard Szeliski. *Computer Vision: Algorithms and Applications*. Texts in Computer Science. London: Springer-Verlag, 2011.


Hoiem, Derek, CS 543 - Computer Vision (Spring 2011). URL: https://courses.engr.illinois.edu/cs543/sp2017/ (visited on 04/25/2020).