

Deep Learning for Computer Vision

# Generative Adversarial Networks

Vineeth N Balasubramanian

Department of Computer Science and Engineering  
Indian Institute of Technology, Hyderabad



## Review: Question

Why does using KL-divergence in finding the generative model simplify to maximum likelihood estimation?

- Recall our problem:

$$\theta^* = \arg \min_{\theta \in M} \text{dist}(p_\theta, p_D)$$

- If distance is KL-divergence:

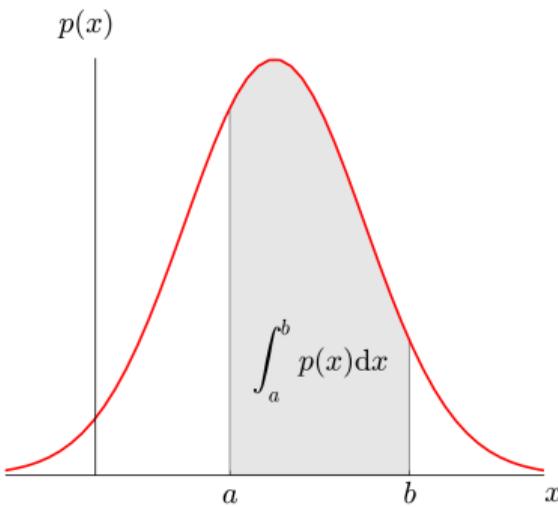
$$\theta^* = \arg \min_{\theta \in M} D_{KL}(p_D || p_\theta)$$

$$= \arg \min_{\theta \in M} \mathbb{E}_{\mathbf{x} \sim p_D} (\log p_D(\mathbf{x}) - \log p_\theta(\mathbf{x}))$$

$$= \arg \min_{\theta \in M} \mathbb{E}_{\mathbf{x} \sim p_D} (-\log p_\theta(\mathbf{x})) \quad (\theta \text{ parameters are independent of } \log p_D(\mathbf{x}))$$

This is maximum likelihood estimation (minimizing negative log likelihood)!

# Recap



- **Generative Models:** Learn probability density function  $p(x)$  over images  $x$
- $p(x)$  assigns positive number to each possible image  $x$   
     $\implies$  how likely is image  $x$  under distribution  $p(x)$
- **Applications:**
  - Sample to generate new data
  - Likelihood estimation(Outlier detection)
  - Feature learning (unsupervised)

# Density Estimation

## Explicit Density Estimation

- Write explicit function  $p(x) = f(x, \theta)$

Input: Image  $x$

Output: Likelihood value for image

Parameter: Weights  $\theta$

- Assign explicit likelihood to images

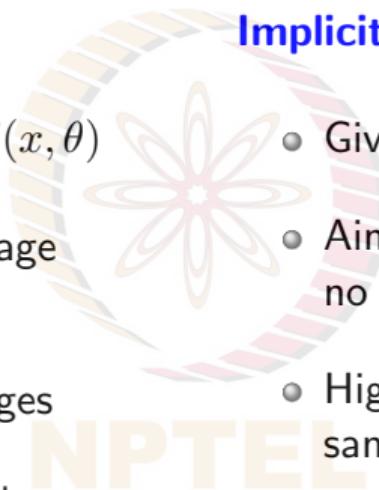
- Can enable outlier detection, but compromise generated image quality/sampling speed

## Implicit Density Estimation

- Give up estimating explicit form of  $p(x)$

- Aim only to sample images from model; no explicit likelihood assignment

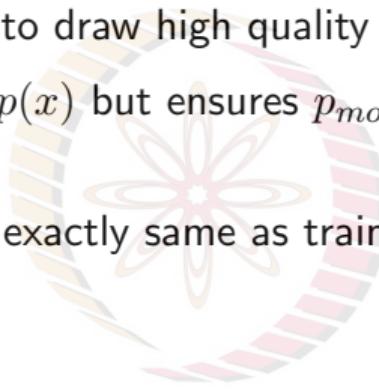
- High quality samples generated/fast sampling speed



# Generative Adversarial Networks (GANs)<sup>1</sup>

## Goal:

- Build good sampler that allows to draw high quality samples from  $p_{model}(x)$
- Not explicitly compute form of  $p(x)$  but ensures  $p_{model}(x)$  approximates/is close to  $p_{data}(x)$
- Output samples similar but not exactly same as train data; how?



NPTEL

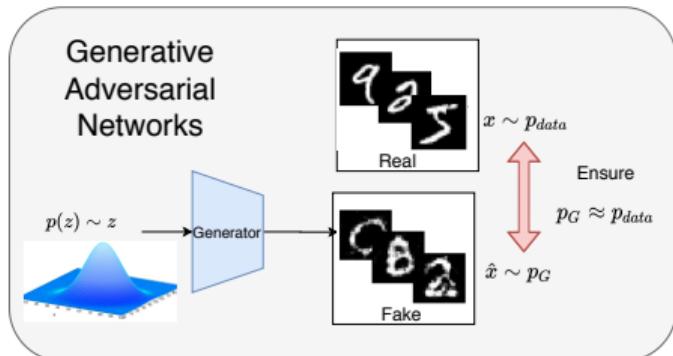
---

<sup>1</sup>Goodfellow et al, Generative Adversarial Nets, NeurIPS 2014

# Generative Adversarial Networks (GANs)<sup>1</sup>

## Goal:

- Build good sampler that allows to draw high quality samples from  $p_{model}(x)$
- Not explicitly compute form of  $p(x)$  but ensures  $p_{model}(x)$  approximates/is close to  $p_{data}(x)$
- Output samples similar but not exactly same as train data; how?



## Method:

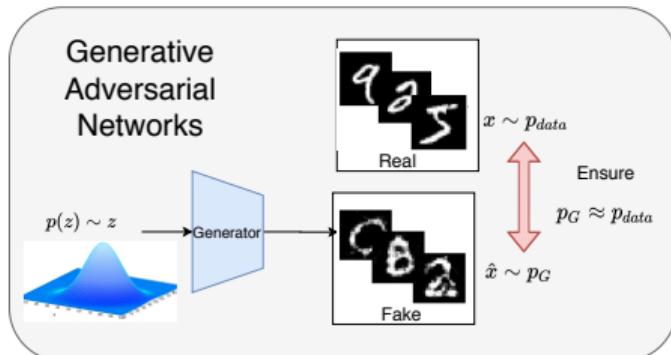
- Introduce latent variable  $z$  with simple prior  $p(z)$  (e.g. Gaussian)

<sup>1</sup>Goodfellow et al, Generative Adversarial Nets, NeurIPS 2014

# Generative Adversarial Networks (GANs)<sup>1</sup>

## Goal:

- Build good sampler that allows to draw high quality samples from  $p_{model}(x)$
- Not explicitly compute form of  $p(x)$  but ensures  $p_{model}(x)$  approximates/is close to  $p_{data}(x)$
- Output samples similar but not exactly same as train data; how?



## Method:

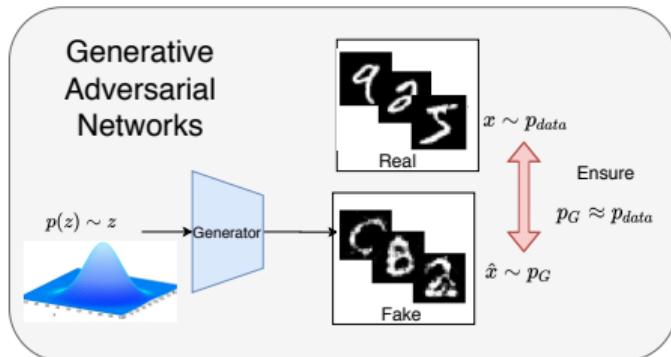
- Introduce latent variable  $z$  with simple prior  $p(z)$  (e.g. Gaussian)
- Sample  $z \sim p(z)$ , pass it through Generator  $\hat{x} = G(z)$ ; where  $\hat{x} \sim p_G$

<sup>1</sup>Goodfellow et al, Generative Adversarial Nets, NeurIPS 2014

# Generative Adversarial Networks (GANs)<sup>1</sup>

## Goal:

- Build good sampler that allows to draw high quality samples from  $p_{model}(x)$
- Not explicitly compute form of  $p(x)$  but ensures  $p_{model}(x)$  approximates/is close to  $p_{data}(x)$
- Output samples similar but not exactly same as train data; how?

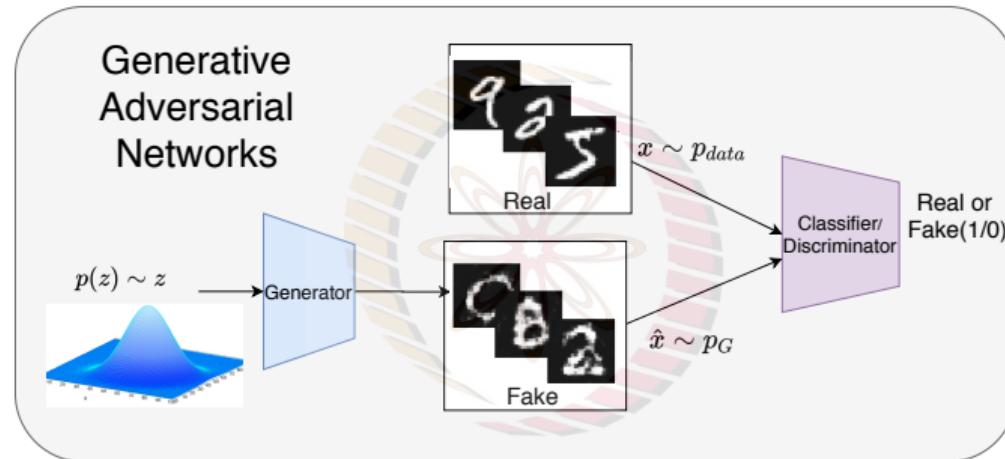


## Method:

- Introduce latent variable  $z$  with simple prior  $p(z)$  (e.g. Gaussian)
- Sample  $z \sim p(z)$ , pass it through Generator  $\hat{x} = G(z)$ ; where  $\hat{x} \sim p_G$
- Introduce mechanism to ensure  $p_G \approx p_{data}$

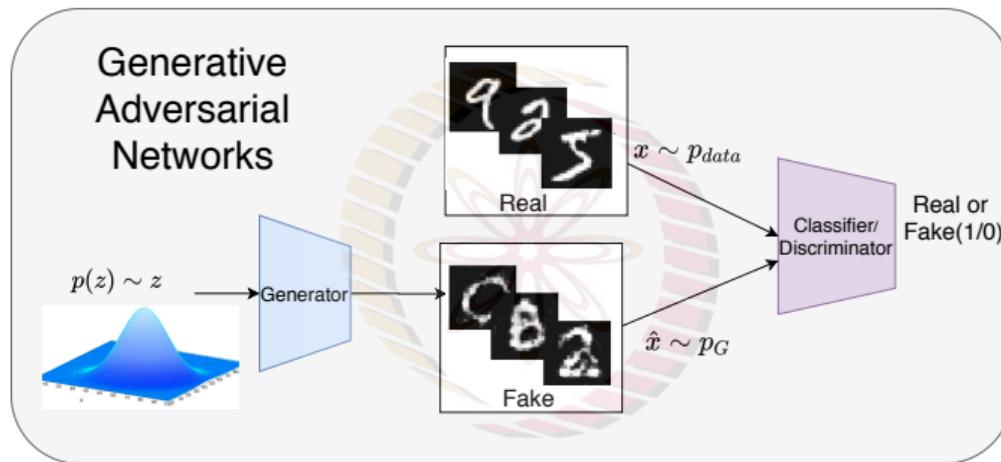
<sup>1</sup>Goodfellow et al, Generative Adversarial Nets, NeurIPS 2014

## GANs: Ensuring $p_G \approx p_{data}$



- Idea: Use a classifier (called **discriminator**) that differentiates between real samples  $x \sim p_{data}$  (class 1) and  $\hat{x} \sim p_G$  (class 0)

## GANs: Ensuring $p_G \approx p_{data}$



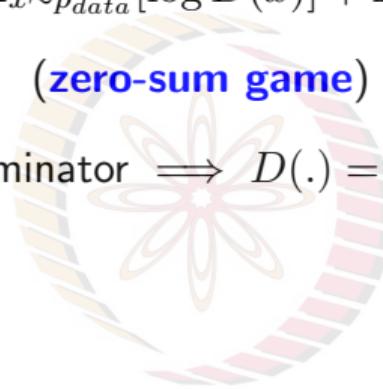
- **Idea:** Use a classifier (called **discriminator**) that differentiates between real samples  $x \sim p_{data}$  (class 1) and  $\hat{x} \sim p_G$  (class 0)
- Train generator  $G$  such that discriminator misclassifies generated sample  $\hat{x}$  into class 1  
     $\Rightarrow$  can no more differentiate between  $x \sim p_{data}$  and  $\hat{x} \sim p_G$

# GANs: How to train?

**Training Objective:**  $\min_G \max_D (\mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] )$

(zero-sum game)

(Assume sigmoid activation in discriminator  $\implies D(\cdot) = \text{probability that input sample is real}$ )



## GANs: How to train?

**Training Objective:**  $\min_G \max_D (\mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] )$

(zero-sum game)

(Assume sigmoid activation in discriminator  $\implies D(\cdot) = \text{probability that input sample is real}$ )

$$O_1 = \max_D \mathbb{E}_{x \sim p_{data}} [\log D(x)]$$

Train discriminator,  $D$ , such that if sample belongs to  $p_{data}$ , maximize the log probability of it being a real sample

## GANs: How to train?

**Training Objective:**  $\min_G \max_D (\mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))] )$

(zero-sum game)

(Assume sigmoid activation in discriminator  $\implies D(\cdot) = \text{probability that input sample is real}$ )

$$O_1 = \max_D \mathbb{E}_{x \sim p_{data}} [\log D(x)]$$

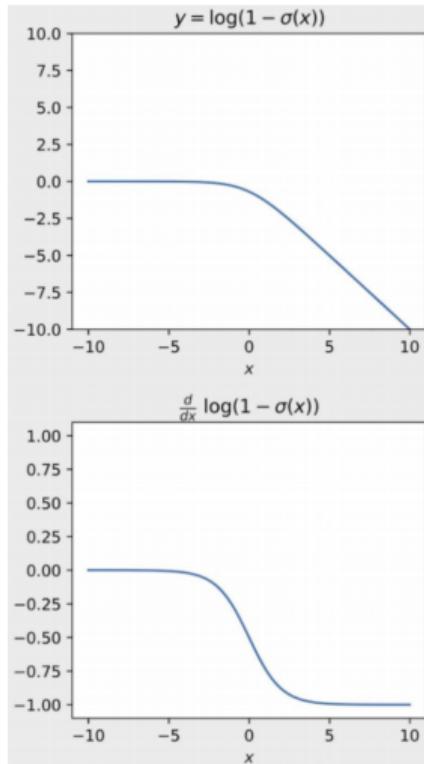
$$O_2 = \min_G \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

Train discriminator,  $D$ , such that if sample belongs to  $p_{data}$ , maximize the log probability of it being a real sample

Train Generator,  $G$  such that if sample belongs to  $p_G$  ( i.e  $G(z)$ ), maximize the log probability of it being a real sample

**Note:** the expectation in the objective function simply implies that losses are averaged over a batch of samples

# GANs: Training Strategy

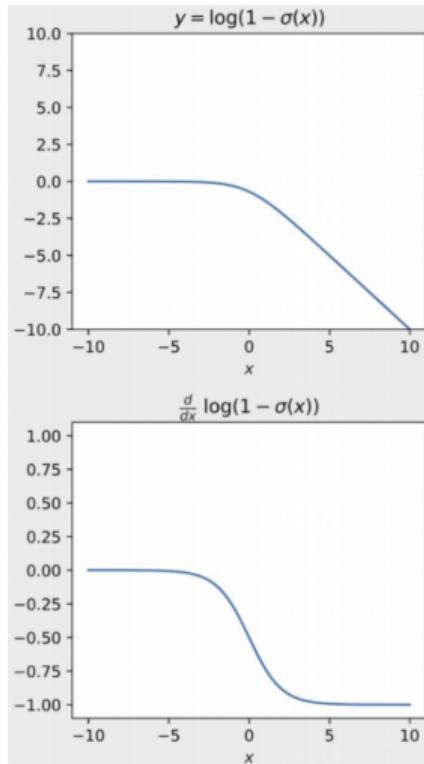


- **Idea:** First train D completely to optimize  $O_1$ ; then train G to optimize  $O_2$



NPTEL

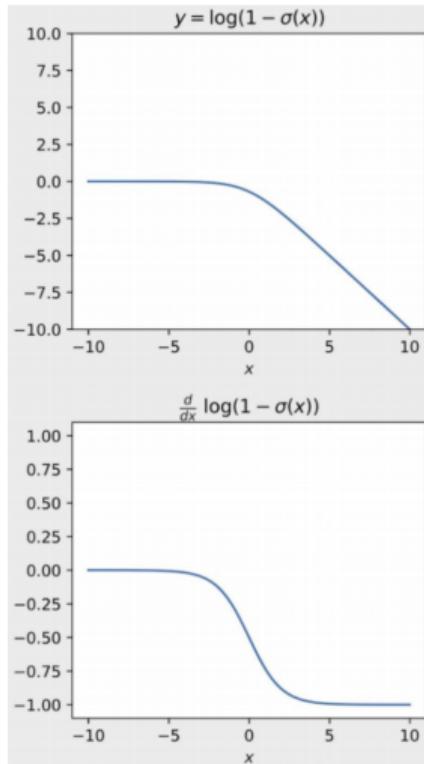
# GANs: Training Strategy



- **Idea:** First train D completely to optimize  $O_1$ ; then train G to optimize  $O_2$
- **Problem:** If D is initially very confident that samples from G are fake, when  $x$  is obtained from G:  
 $D(x)$  or  $\sigma(x) = 0 \implies \log(1 - \sigma(x)) = 0$

NPTEL

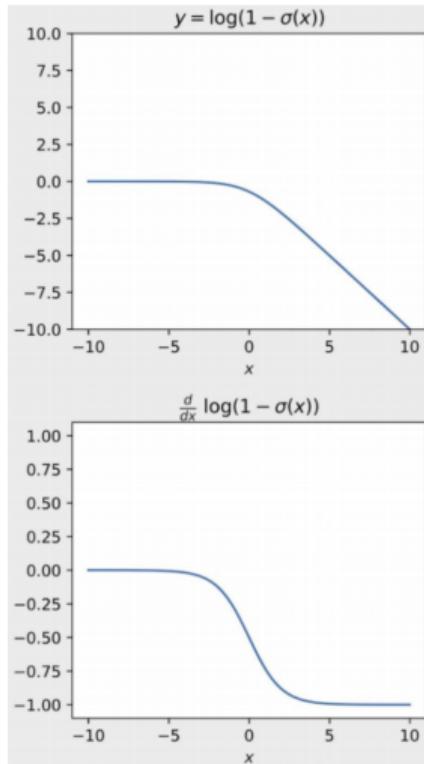
# GANs: Training Strategy



- **Idea:** First train D completely to optimize  $O_1$ ; then train G to optimize  $O_2$
- **Problem:** If D is initially very confident that samples from G are fake, when  $x$  is obtained from G:  
 $D(x)$  or  $\sigma(x) = 0 \implies \log(1 - \sigma(x)) = 0$   
 $\implies \frac{\partial(\log(1-\sigma(x)))}{\partial x} = 0$  (as shown in graph)

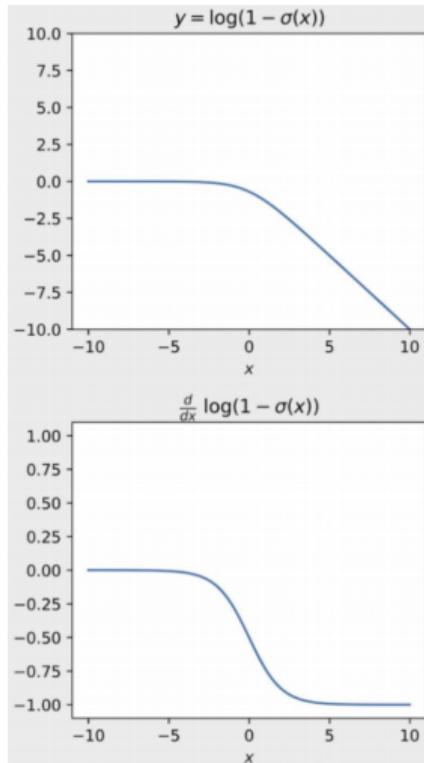
NPTEL

# GANs: Training Strategy



- **Idea:** First train D completely to optimize  $O_1$ ; then train G to optimize  $O_2$
- **Problem:** If D is initially very confident that samples from G are fake, when  $x$  is obtained from G:  
 $D(x)$  or  $\sigma(x) = 0 \implies \log(1 - \sigma(x)) = 0$   
 $\implies \frac{\partial(\log(1-\sigma(x)))}{\partial x} = 0$  (as shown in graph)  
 $\implies$  G gets no gradients to train!

# GANs: Training Strategy



- **Idea:** First train D completely to optimize  $O_1$ ; then train G to optimize  $O_2$
- **Problem:** If D is initially very confident that samples from G are fake, when  $x$  is obtained from G:  
 $D(x)$  or  $\sigma(x) = 0 \implies \log(1 - \sigma(x)) = 0$   
 $\implies \frac{\partial(\log(1-\sigma(x)))}{\partial x} = 0$  (as shown in graph)  
 $\implies$  G gets no gradients to train!
- **Solution:** Alternate between Discriminator objective  $O_1$  and Generator objective  $O_2$

# GANs: Algorithm<sup>2</sup>

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

**end for**

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

<sup>2</sup>Goodfellow et al, Generative Adversarial Nets, NeurIPS 2014

# GANs: Evaluating Optimality

**Objective:**  $\min_G \max_D \quad \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$



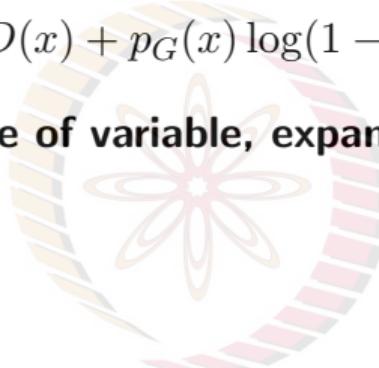
NPTEL

## GANs: Evaluating Optimality

**Objective:**  $\min_G \max_D \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$

$$\implies \min_G \max_D \int_x (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx$$

(change of variable, expanding expectation)



## GANs: Evaluating Optimality

**Objective:**  $\min_G \max_D \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$

$$\implies \min_G \max_D \int_x (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx$$

(change of variable, expanding expectation)

$$\implies \min_G \int_x \max_D (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx \quad (\text{taking max inside})$$

NPTEL

## GANs: Evaluating Optimality

**Objective:**  $\min_G \max_D \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$

$$\implies \min_G \max_D \int_x (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx$$

(change of variable, expanding expectation)

$$\implies \min_G \int_x \max_D (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx \quad (\text{taking max inside})$$

Let  $y = D(x); a = p_{data}; b = p_G$

NPTEL

## GANs: Evaluating Optimality

**Objective:**  $\min_G \max_D \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$

$$\implies \min_G \max_D \int_x (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx$$

**(change of variable, expanding expectation)**

$$\implies \min_G \int_x \max_D (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx \quad \text{(taking max inside)}$$

Let  $y = D(x); a = p_{data}; b = p_G$

Then,  $f(y) = a \log y + b \log(1 - y)$

## GANs: Evaluating Optimality

**Objective:**  $\min_G \max_D \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$

$$\implies \min_G \max_D \int_x (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx$$

**(change of variable, expanding expectation)**

$$\implies \min_G \int_x \max_D (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx \quad \text{(taking max inside)}$$

Let  $y = D(x); a = p_{data}; b = p_G$

Then,  $f(y) = a \log y + b \log(1 - y)$

$$f'(y) = \frac{a}{y} - \frac{b}{1-y}; f'(y) = 0 \implies y = \frac{a}{a+b} \quad \text{(local max)}$$

## GANs: Evaluating Optimality

**Objective:**  $\min_G \max_D \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$

$$\implies \min_G \max_D \int_x (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx$$

**(change of variable, expanding expectation)**

$$\implies \min_G \int_x \max_D (p_{data}(x) \log D(x) + p_G(x) \log(1 - D(x))) dx \quad \text{(taking max inside)}$$

Let  $y = D(x); a = p_{data}; b = p_G$

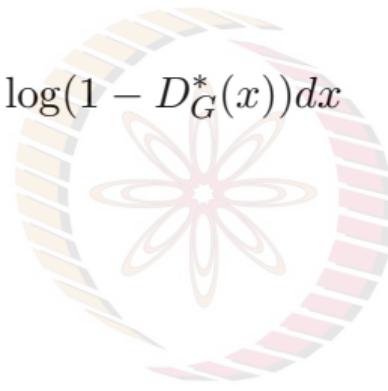
Then,  $f(y) = a \log y + b \log(1 - y)$

$$f'(y) = \frac{a}{y} - \frac{b}{1-y}; f'(y) = 0 \implies y = \frac{a}{a+b} \quad \text{(local max)}$$

**Optimal Discriminator:**  $D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}$

# GANs: Evaluating Optimality

$$\min_G \int_X (p_{data}(x) \log D_G^*(x) + p_G(x) \log(1 - D_G^*(x))) dx$$

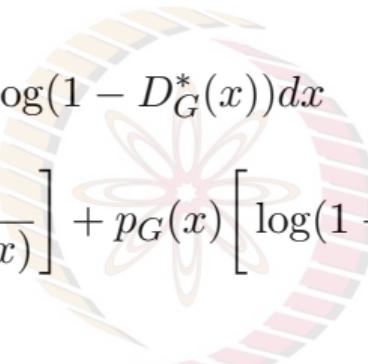


NPTEL

# GANs: Evaluating Optimality

$$\min_G \int_X (p_{data}(x) \log D_G^*(x) + p_G(x) \log(1 - D_G^*(x))) dx$$

$$\min_G \int_X (p_{data}(x) \left[ \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + p_G(x) \left[ \log(1 - \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}) \right]) dx$$



NPTEL

## GANs: Evaluating Optimality

$$\min_G \int_X (p_{data}(x) \log D_G^*(x) + p_G(x) \log(1 - D_G^*(x))) dx$$

$$\min_G \int_X (p_{data}(x) \left[ \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + p_G(x) \left[ \log(1 - \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}) \right]) dx$$

$$\min_G \int_X (p_{data}(x) \left[ \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + p_G(x) \left[ \log(\frac{p_G(x)}{p_{data}(x) + p_G(x)}) \right]) dx$$

NPTEL

## GANs: Evaluating Optimality

$$\min_G \int_X (p_{data}(x) \log D_G^*(x) + p_G(x) \log(1 - D_G^*(x))) dx$$

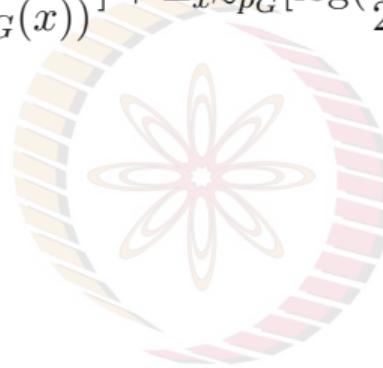
$$\min_G \int_X (p_{data}(x) \left[ \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + p_G(x) \left[ \log(1 - \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}) \right]) dx$$

$$\min_G \int_X (p_{data}(x) \left[ \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + p_G(x) \left[ \log(\frac{p_G(x)}{p_{data}(x) + p_G(x)}) \right]) dx$$

$$\min_G \left( \mathbb{E}_{x \sim p_{data}} [\log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}] + \mathbb{E}_{x \sim p_G} [\log(\frac{p_G(x)}{p_{data}(x) + p_G(x)})] \right)$$

# GANs: Evaluating Optimality

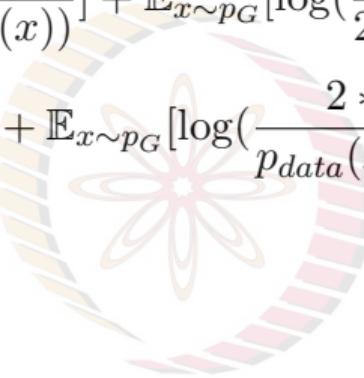
$$\min_G \left( \mathbb{E}_{x \sim p_{data}} [\log \frac{2 * p_{data}(x)}{2 * (p_{data}(x) + p_G(x))}] + \mathbb{E}_{x \sim p_G} [\log (\frac{2 * p_G(x)}{2 * (p_{data}(x) + p_G(x))})] \right)$$



NPTEL

# GANs: Evaluating Optimality

$$\min_G \left( \mathbb{E}_{x \sim p_{data}} [\log \frac{2 * p_{data}(x)}{2 * (p_{data}(x) + p_G(x))}] + \mathbb{E}_{x \sim p_G} [\log(\frac{2 * p_G(x)}{2 * (p_{data}(x) + p_G(x))})] \right)$$
$$\min_G \left( \mathbb{E}_{x \sim p_{data}} [\log \frac{2 * p_{data}(x)}{p_{data}(x) + p_G(x)}] + \mathbb{E}_{x \sim p_G} [\log(\frac{2 * p_G(x)}{p_{data}(x) + p_G(x)})] - \log 4 \right)$$



NPTEL

## GANs: Evaluating Optimality

$$\min_G \left( \mathbb{E}_{x \sim p_{data}} [\log \frac{2 * p_{data}(x)}{2 * (p_{data}(x) + p_G(x))}] + \mathbb{E}_{x \sim p_G} [\log(\frac{2 * p_G(x)}{2 * (p_{data}(x) + p_G(x))})] \right)$$
$$\min_G \left( \mathbb{E}_{x \sim p_{data}} [\log \frac{2 * p_{data}(x)}{p_{data}(x) + p_G(x)}] + \mathbb{E}_{x \sim p_G} [\log(\frac{2 * p_G(x)}{p_{data}(x) + p_G(x)})] - \log 4 \right)$$
$$\min_G \left( \mathbf{KL}(p_{data}(x), \frac{p_{data}(x) + p_G(x)}{2}) + \mathbf{KL}(p_G(x), \frac{p_{data}(x) + p_G(x)}{2}) - \log 4 \right)$$

NPTEL

## GANs: Evaluating Optimality

$$\begin{aligned} & \min_G \left( \mathbb{E}_{x \sim p_{data}} [\log \frac{2 * p_{data}(x)}{2 * (p_{data}(x) + p_G(x))}] + \mathbb{E}_{x \sim p_G} [\log (\frac{2 * p_G(x)}{2 * (p_{data}(x) + p_G(x))})] \right) \\ & \min_G \left( \mathbb{E}_{x \sim p_{data}} [\log \frac{2 * p_{data}(x)}{p_{data}(x) + p_G(x)}] + \mathbb{E}_{x \sim p_G} [\log (\frac{2 * p_G(x)}{p_{data}(x) + p_G(x)})] - \log 4 \right) \\ & \min_G \left( \mathbf{KL}(p_{data}(x), \frac{p_{data}(x) + p_G(x)}{2}) + \mathbf{KL}(p_G(x), \frac{p_{data}(x) + p_G(x)}{2}) - \log 4 \right) \end{aligned}$$

Kullback–Leibler Divergence

$$\mathbf{KL}(\mathbf{p}, \mathbf{q}) = \mathbb{E}_{x \sim p} [\frac{\log p(x)}{\log q(x)}]$$

Jenson-Shannon Divergence

$$\mathbf{JSD}(p_{data}, p_G) = \frac{KL(p_{data}, \frac{p_{data}+p_G}{2})}{2} + \frac{KL(p_G, \frac{p_{data}+p_G}{2})}{2}$$

$\min_G (2 * JSD(p_{data}, p_G) - \log 4) \implies$  expression is minimized when  $p_{data} = p_G$

(Recall  $JSD \geq 0$  by definition)

# Optimality of GANs: Conclusions

- $D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x)+p_G(x)}$  (Optimal Discriminator for any G)



NPTEL

# Optimality of GANs: Conclusions

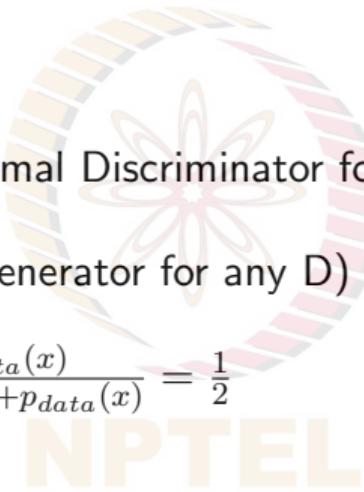
- $D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x)+p_G(x)}$  (Optimal Discriminator for any G)
- $p_{data}(x) = p_G(x)$  (Optimal Generator for any D)



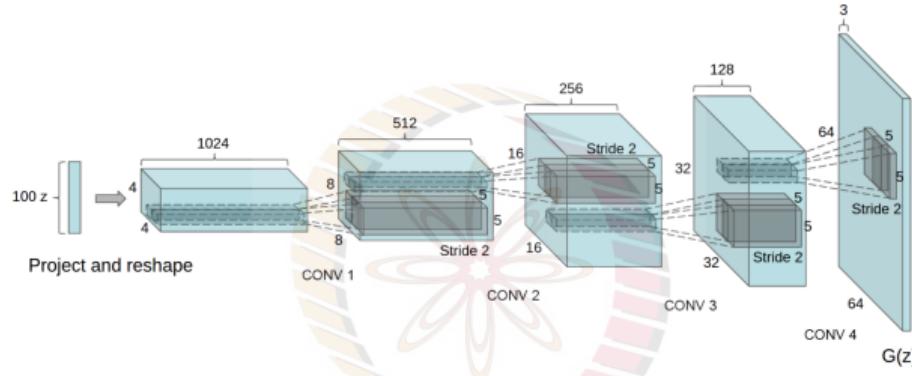
NPTEL

# Optimality of GANs: Conclusions

- $D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x)+p_G(x)}$  (Optimal Discriminator for any G)
- $p_{data}(x) = p_G(x)$  (Optimal Generator for any D)
- $D_G^*(x) = \frac{p_G(x)}{p_G(x)+p_G(x)} = \frac{p_{data}(x)}{p_{data}(x)+p_{data}(x)} = \frac{1}{2}$



# Deep Convolution GAN (DCGAN)<sup>3</sup>

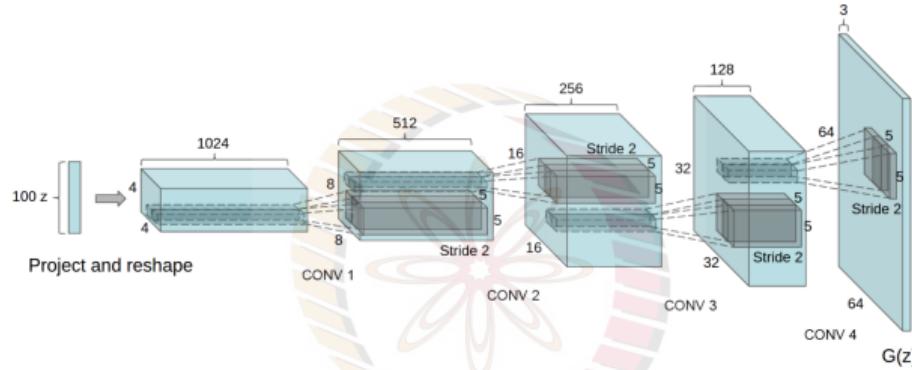


- Bridge the gap between success of CNNs for supervised learning and unsupervised generative models

NPTEL

<sup>3</sup>Radford et al, Unsupervised Representation learning with deep convolutional GANs, NeurIPS 2016

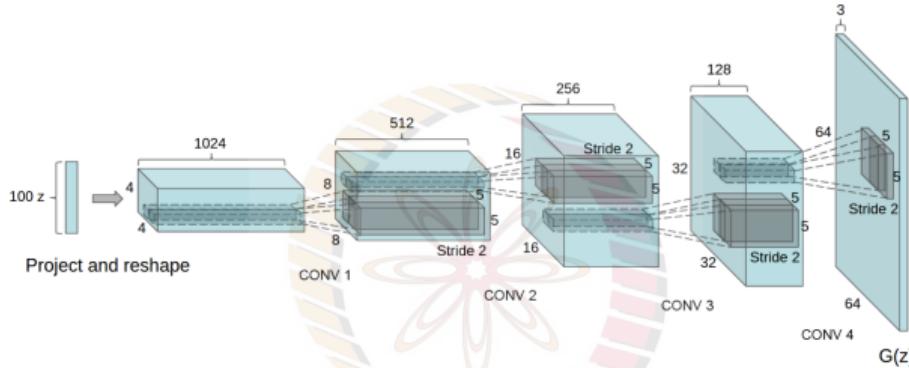
# Deep Convolution GAN (DCGAN)<sup>3</sup>



- Bridge the gap between success of CNNs for supervised learning and unsupervised generative models
- Best practices to enable stable training of GANs with deep architectures

<sup>3</sup>Radford et al, Unsupervised Representation learning with deep convolutional GANs, NeurIPS 2016

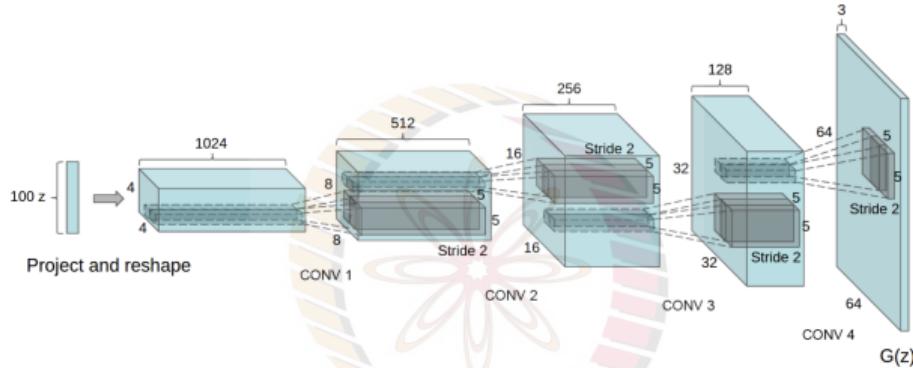
# Deep Convolution GAN (DCGAN)<sup>3</sup>



- Bridge the gap between success of CNNs for supervised learning and unsupervised generative models
- Best practices to enable stable training of GANs with deep architectures
- Smooth high-dimensional Interpolations; Vector arithmetic

<sup>3</sup>Radford et al, Unsupervised Representation learning with deep convolutional GANs, NeurIPS 2016

# Deep Convolution GAN (DCGAN)<sup>3</sup>



- Bridge the gap between success of CNNs for supervised learning and unsupervised generative models
- Best practices to enable stable training of GANs with deep architectures
- Smooth high-dimensional Interpolations; Vector arithmetic
- Demonstrate unsupervised feature learning capabilities of GANs and its applications  
**(Representation Learning)**

<sup>3</sup>Radford et al, Unsupervised Representation learning with deep convolutional GANs, NeurIPS 2016

# DCGAN: Training Practices for Deeper GANs<sup>4</sup>

- Replaces deterministic spatial pooling functions (e.g maxpooling) with **strided convolutions** → allows to learn spatial downsampling

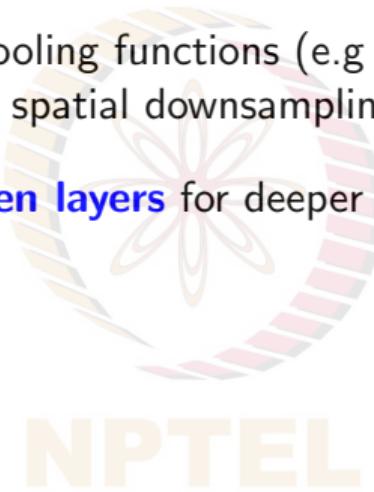


---

<sup>4</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

# DCGAN: Training Practices for Deeper GANs<sup>4</sup>

- Replaces deterministic spatial pooling functions (e.g maxpooling) with **strided convolutions** → allows to learn spatial downsampling
- **Remove fully connected hidden layers** for deeper architectures



---

<sup>4</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

# DCGAN: Training Practices for Deeper GANs<sup>4</sup>

- Replaces deterministic spatial pooling functions (e.g maxpooling) with **strided convolutions** → allows to learn spatial downsampling
- **Remove fully connected hidden layers** for deeper architectures
- **Batchnorm in G and D** → prevent generator collapse/enable gradient flow in deeper architectures; not applied at output of G and input of D

NPTEL

---

<sup>4</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

# DCGAN: Training Practices for Deeper GANs<sup>4</sup>

- Replaces deterministic spatial pooling functions (e.g maxpooling) with **strided convolutions** → allows to learn spatial downsampling
- **Remove fully connected hidden layers** for deeper architectures
- **Batchnorm in G and D** → prevent generator collapse/enable gradient flow in deeper architectures; not applied at output of G and input of D
- **Non-Linearity:** ReLU for generator, Leaky-ReLU (0.2) for discriminator

<sup>4</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

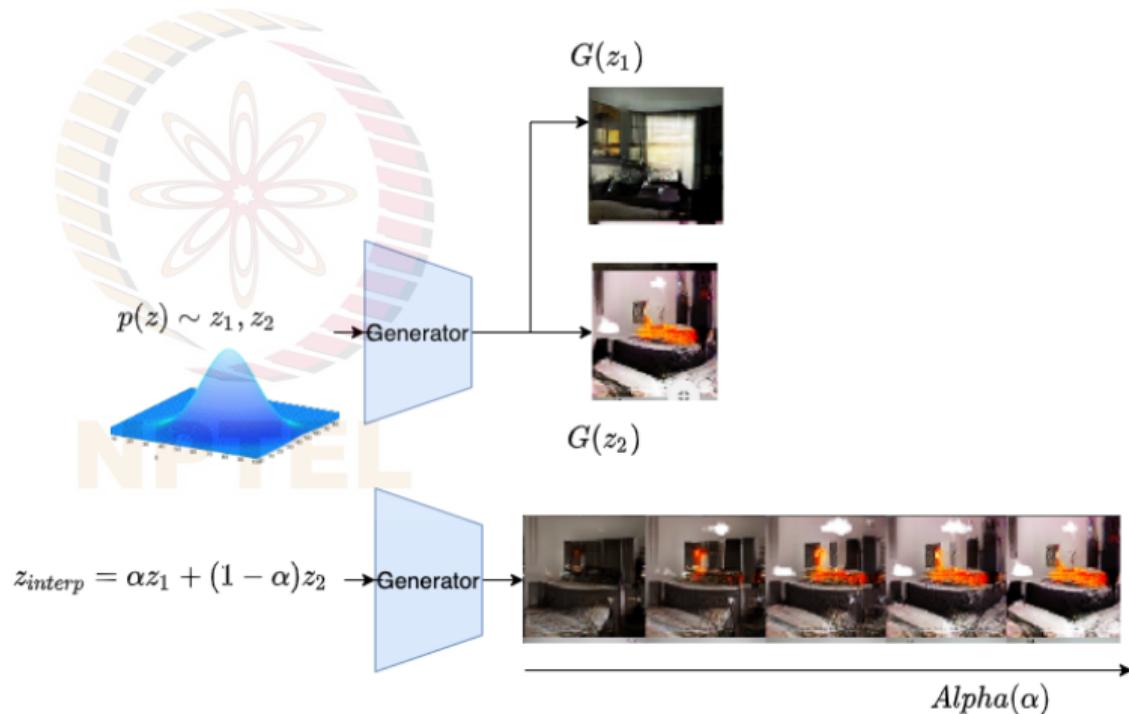
# DCGAN: Training Practices for Deeper GANs<sup>4</sup>

- Replaces deterministic spatial pooling functions (e.g maxpooling) with **strided convolutions** → allows to learn spatial downsampling
- **Remove fully connected hidden layers** for deeper architectures
- **Batchnorm in G and D** → prevent generator collapse/enable gradient flow in deeper architectures; not applied at output of G and input of D
- **Non-Linearity:** ReLU for generator, Leaky-ReLU (0.2) for discriminator
- **Output Non-Linearity:** tanh for generator, sigmoid for discriminator

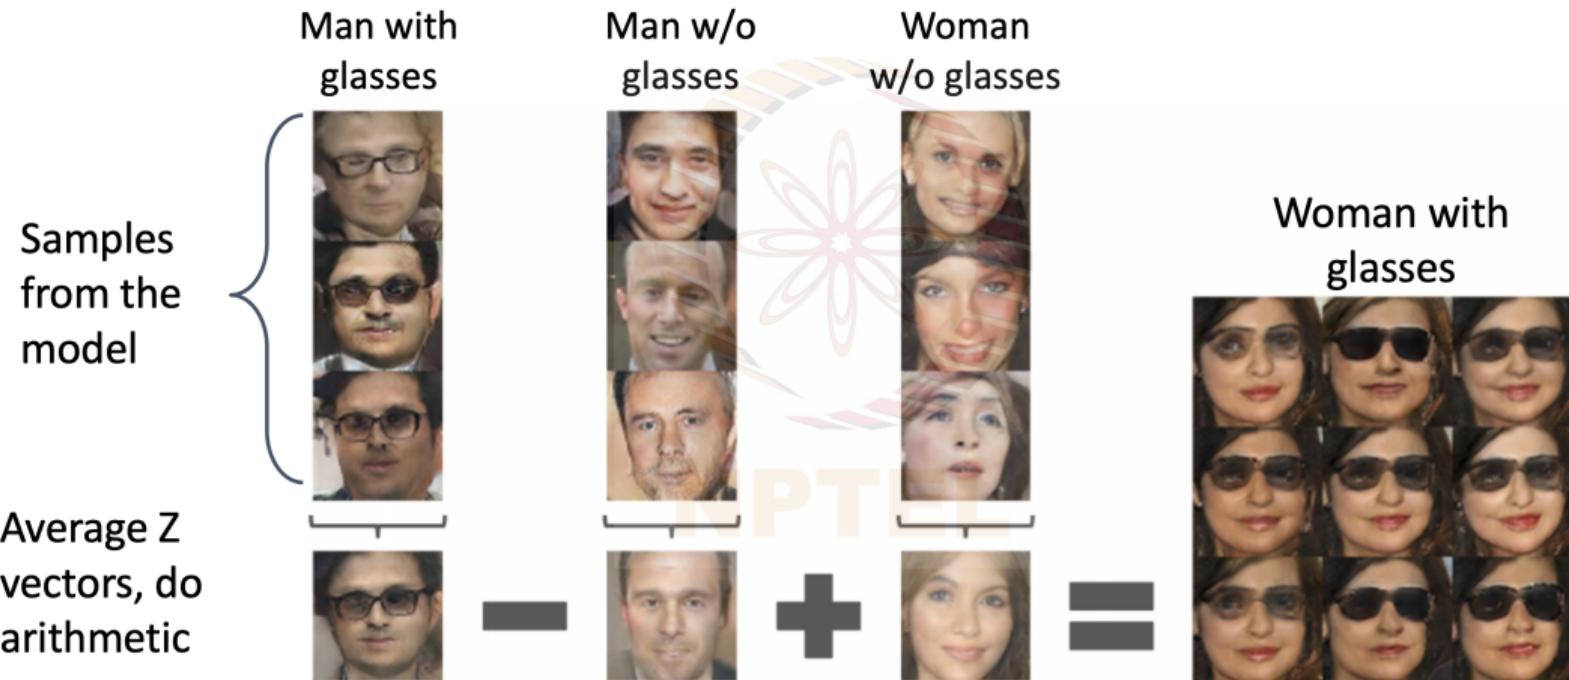
<sup>4</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

# Walking the Latent Space

- We can interpolate between two points in latent space and visualize
- Smooth transition in generated image space by changing latent (perceptual) features is a sign of a good model



## Vector Arithmetic in Latent Space<sup>5</sup>



<sup>5</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

# Pose Transformation<sup>6</sup>

- $z_{\text{left}}$ : averaged latent vector of faces looking left  
 $z_{\text{right}}$ : average latent vector of faces looking right
- $z_{\text{turn}} = z_{\text{right}} - z_{\text{left}}$   
 $z_{\text{new}} = z + \alpha z_{\text{turn}}$
- $G(z_{\text{new}}) \implies \text{Transformed image}$



<sup>6</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

# Feature Learning<sup>7</sup>

Table 1: CIFAR-10 classification results using our pre-trained model. Our DCGAN is not pre-trained on CIFAR-10, but on Imagenet-1k, and the features are used to classify CIFAR-10 images.

Model	Accuracy	Accuracy (400 per class)	max # of features units
1 Layer K-means	80.6%	63.7% ( $\pm 0.7\%$ )	4800
3 Layer K-means Learned RF	82.0%	70.7% ( $\pm 0.7\%$ )	3200
View Invariant K-means	81.9%	72.6% ( $\pm 0.7\%$ )	6400
Exemplar CNN	84.3%	77.4% ( $\pm 0.2\%$ )	1024
DCGAN (ours) + L2-SVM	82.8%	73.8% ( $\pm 0.4\%$ )	512

- Train DCGAN on ImageNet-1k

NPTEL

<sup>7</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

# Feature Learning<sup>7</sup>

Table 1: CIFAR-10 classification results using our pre-trained model. Our DCGAN is not pre-trained on CIFAR-10, but on Imagenet-1k, and the features are used to classify CIFAR-10 images.

Model	Accuracy	Accuracy (400 per class)	max # of features units
1 Layer K-means	80.6%	63.7% ( $\pm 0.7\%$ )	4800
3 Layer K-means Learned RF	82.0%	70.7% ( $\pm 0.7\%$ )	3200
View Invariant K-means	81.9%	72.6% ( $\pm 0.7\%$ )	6400
Exemplar CNN	84.3%	77.4% ( $\pm 0.2\%$ )	1024
DCGAN (ours) + L2-SVM	82.8%	73.8% ( $\pm 0.4\%$ )	512

- Train DCGAN on ImageNet-1k
- Use discriminator's convolution features for supervised tasks i.e. classify images from CIFAR-10 dataset

NPTEL

<sup>7</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

# Feature Learning<sup>7</sup>

Table 1: CIFAR-10 classification results using our pre-trained model. Our DCGAN is not pre-trained on CIFAR-10, but on Imagenet-1k, and the features are used to classify CIFAR-10 images.

Model	Accuracy	Accuracy (400 per class)	max # of features units
1 Layer K-means	80.6%	63.7% ( $\pm 0.7\%$ )	4800
3 Layer K-means Learned RF	82.0%	70.7% ( $\pm 0.7\%$ )	3200
View Invariant K-means	81.9%	72.6% ( $\pm 0.7\%$ )	6400
Exemplar CNN	84.3%	77.4% ( $\pm 0.2\%$ )	1024
DCGAN (ours) + L2-SVM	82.8%	73.8% ( $\pm 0.4\%$ )	512

- Train DCGAN on ImageNet-1k
- Use discriminator's convolution features for supervised tasks i.e. classify images from CIFAR-10 dataset
- DCGAN never trained on CIFAR-10  $\implies$  representation power and domain robustness of learned features

<sup>7</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

# Feature Learning<sup>7</sup>

Table 1: CIFAR-10 classification results using our pre-trained model. Our DCGAN is not pre-trained on CIFAR-10, but on Imagenet-1k, and the features are used to classify CIFAR-10 images.

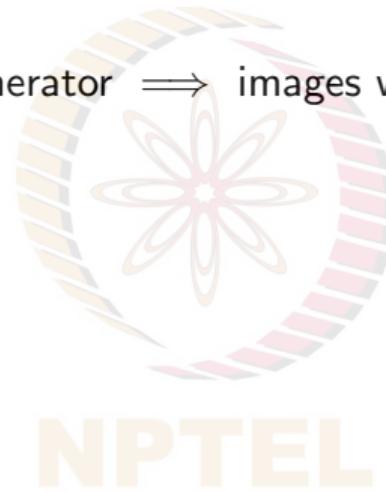
Model	Accuracy	Accuracy (400 per class)	max # of features units
1 Layer K-means	80.6%	63.7% ( $\pm 0.7\%$ )	4800
3 Layer K-means Learned RF	82.0%	70.7% ( $\pm 0.7\%$ )	3200
View Invariant K-means	81.9%	72.6% ( $\pm 0.7\%$ )	6400
Exemplar CNN	84.3%	77.4% ( $\pm 0.2\%$ )	1024
DCGAN (ours) + L2-SVM	82.8%	73.8% ( $\pm 0.4\%$ )	512

- Train DCGAN on ImageNet-1k
- Use discriminator's convolution features for supervised tasks i.e. classify images from CIFAR-10 dataset
- DCGAN never trained on CIFAR-10  $\implies$  representation power and domain robustness of learned features
- Competitive results when compared with other feature learning methods

<sup>7</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

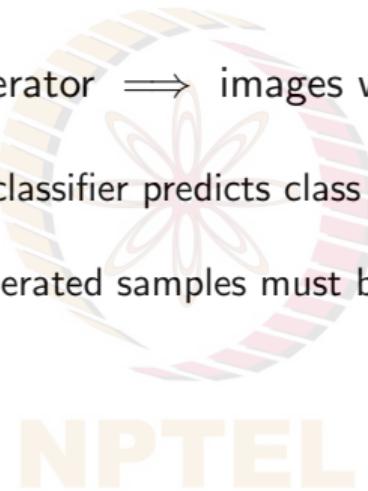
# GANs: How to evaluate?

- **Human judgement:** Good generator  $\implies$  images with distinctly recognizable objects;  
Semantically diverse samples



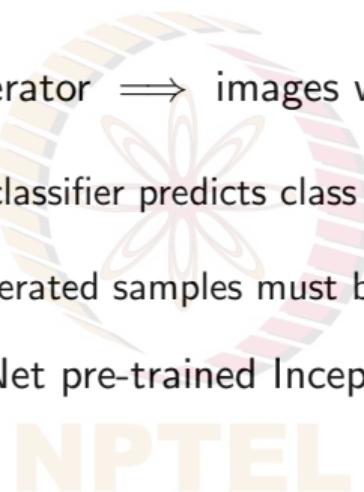
# GANs: How to evaluate?

- **Human judgement:** Good generator  $\implies$  images with distinctly recognizable objects; Semantically diverse samples
  - **Recognizable objects**  $\implies$  classifier predicts class of generated sample correctly with high confidence
  - **Semantic diversity**  $\implies$  generated samples must be evenly from all classes in train set



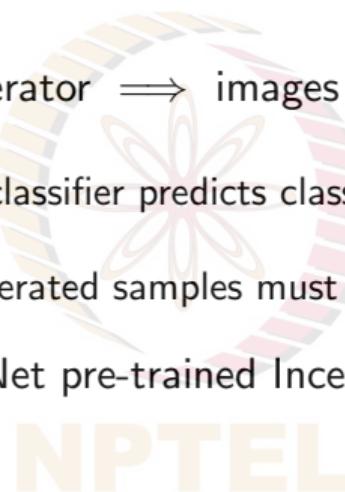
## GANs: How to evaluate?

- **Human judgement:** Good generator  $\implies$  images with distinctly recognizable objects; Semantically diverse samples
  - **Recognizable objects**  $\implies$  classifier predicts class of generated sample correctly with high confidence
  - **Semantic diversity**  $\implies$  generated samples must be evenly from all classes in train set
- **Prediction power:** How ImageNet pre-trained Inception Network V3 performs on these generated images



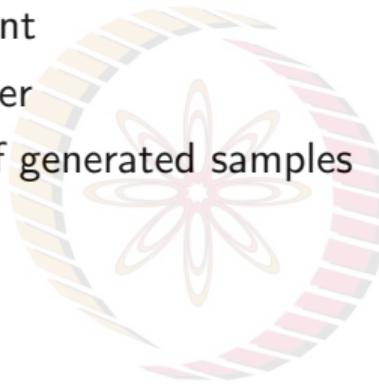
# GANs: How to evaluate?

- **Human judgement:** Good generator  $\implies$  images with distinctly recognizable objects; Semantically diverse samples
  - **Recognizable objects**  $\implies$  classifier predicts class of generated sample correctly with high confidence
  - **Semantic diversity**  $\implies$  generated samples must be evenly from all classes in train set
- **Prediction power:** How ImageNet pre-trained Inception Network V3 performs on these generated images
- Still an open research problem...



# Inception Score<sup>8</sup>

- Correlates with human judgement
- $p(y|x)$ : Inception model/classifier
- $p(y)$ : Label/class distribution of generated samples



<sup>8</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

## Inception Score<sup>8</sup>

- Correlates with human judgement
- $p(y|x)$ : Inception model/classifier
- $p(y)$ : Label/class distribution of generated samples
- **Inception Score (IS):**
$$\begin{aligned} &= \exp(\mathbb{E}_{x \sim p_g}[D_{KL}[p(y|x) || p(y)]]) \\ &= \exp(\mathbb{E}_{x \sim p_g, y \sim p(y|x)}[\log(p(y|x)) - \log(p(y))]) \\ &= \exp(H(y) - H(y|x)) \end{aligned}$$

$H(y)$ : entropy of generated sample class labels (**Semantic diversity = high  $H(y)$** )

$H(y|x)$ : entropy of class labels from classifier (**Distinctly Classifiable = low  $H(y|x)$** )

---

<sup>8</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

## Inception Score<sup>8</sup>

- Correlates with human judgement
- $p(y|x)$ : Inception model/classifier
- $p(y)$ : Label/class distribution of generated samples
- **Inception Score (IS):**
$$\begin{aligned} &= \exp(\mathbb{E}_{x \sim p_g}[D_{KL}[p(y|x) || p(y)]]) \\ &= \exp(\mathbb{E}_{x \sim p_g, y \sim p(y|x)}[\log(p(y|x)) - \log(p(y))]) \\ &= \exp(H(y) - H(y|x)) \end{aligned}$$

$H(y)$ : entropy of generated sample class labels (**Semantic diversity = high  $H(y)$** )

$H(y|x)$ : entropy of class labels from classifier (**Distinctly Classifiable = low  $H(y|x)$** )

- Higher IS  $\implies$  better generative model

---

<sup>8</sup>Radford et al, Unsupervised Representation Learning with Deep Convolutional GANs, NeurIPS 2016

## Fréchet Inception Distance (FID)<sup>9</sup>

- **Drawback of IS:** Statistics of real-world samples are not used to compare with statistics of synthetic samples
- FID enables us to capture more subtle differences; measures diversity better

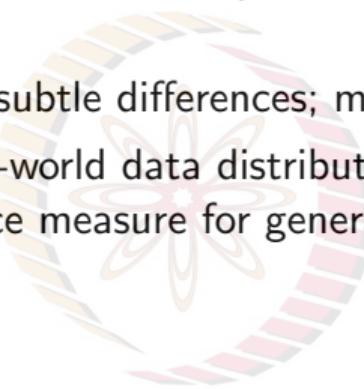


---

<sup>9</sup>Heusel et al, GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium, NeurIPS 2017

## Fréchet Inception Distance (FID)<sup>9</sup>

- **Drawback of IS:** Statistics of real-world samples are not used to compare with statistics of synthetic samples
- FID enables us to capture more subtle differences; measures diversity better
- **Intuition:** Distance between real-world data distribution and generating model's data distribution serves as performance measure for generative models; how?



---

<sup>9</sup>Heusel et al, GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium, NeurIPS 2017

## Fréchet Inception Distance (FID)<sup>9</sup>

- **Drawback of IS:** Statistics of real-world samples are not used to compare with statistics of synthetic samples
- FID enables us to capture more subtle differences; measures diversity better
- **Intuition:** Distance between real-world data distribution and generating model's data distribution serves as performance measure for generative models; how?
- Embed image  $x$  into feature space (activations of Inception-v3 pool3 layer), and compute the Fréchet distance between two multivariate Gaussians:

$$d^2((m, C), (m_w, C_w)) = \|m - m_w\|_2^2 + \text{Tr}(C + C_w - 2(C \cdot C_w)^2)$$

where  $m$  is mean,  $C$  is covariance of generated image features;  $m_w$  is mean,  $C_w$  is covariance of original image features

<sup>9</sup>Heusel et al, GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium, NeurIPS 2017

## Fréchet Inception Distance (FID)<sup>9</sup>

- **Drawback of IS:** Statistics of real-world samples are not used to compare with statistics of synthetic samples
- FID enables us to capture more subtle differences; measures diversity better
- **Intuition:** Distance between real-world data distribution and generating model's data distribution serves as performance measure for generative models; how?
- Embed image  $x$  into feature space (activations of Inception-v3 pool3 layer), and compute the Fréchet distance between two multivariate Gaussians:

$$d^2((m, C), (m_w, C_w)) = \|m - m_w\|_2^2 + \text{Tr}(C + C_w - 2(C \cdot C_w)^2)$$

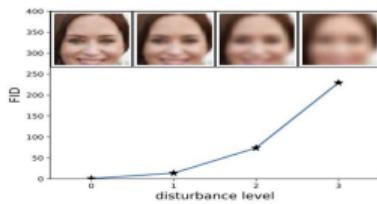
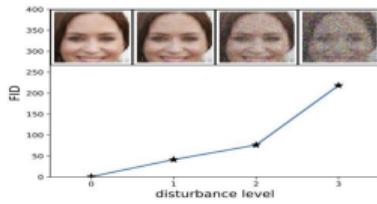
where  $m$  is mean,  $C$  is covariance of generated image features;  $m_w$  is mean,  $C_w$  is covariance of original image features

- Lower FID  $\implies$  better generative model

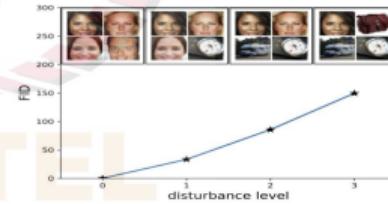
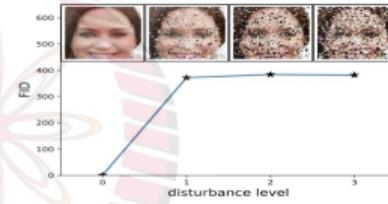
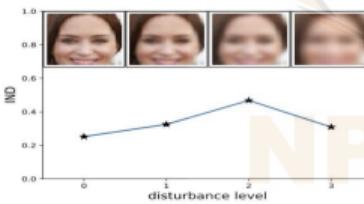
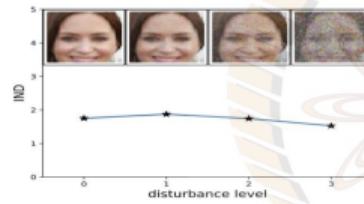
<sup>9</sup>Heusel et al, GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium, NeurIPS 2017

# FID vs Inception Score<sup>10</sup>

CelebA dataset (Columns 1,3: FID Score; Columns 2,4: Inception Score)



(Top:) Gaussian noise added to images  
(Bottom:) Gaussian blur added to images



(Top:) Salt-and-pepper noise added to images  
(Bottom:) ImageNet crops added to images

<sup>10</sup>Heusel et al, GANs Trained by a Two Time-Scale Update Rule Converge to a Local Nash Equilibrium, NeurIPS 2017

# Homework

## Readings

- Dive into Deep Learning, Chapter 17  
[https://d2l.ai/chapter\\_generative-adversarial-networks/gan.html](https://d2l.ai/chapter_generative-adversarial-networks/gan.html)
- Play with Generative Adversarial Networks (GANs) in your browser  
<https://poloclub.github.io/ganlab/>
- Deep Learning With Python - François Chollet, Section 8.5  
<https://github.com/veritone/ds-transcription-capstone>